

Introdução

O código visa aproximar a equação de laplace utilizando o método Gauss-Seidel, que é um método iterativo para resolução de sistema de equações lineares, tendo a iteração definida por:

$$x^{(k+1)} = (D + L)^{-1} \left(-Ux^{(k)} + b \right),$$

Neste caso, a perfilagem do código tem grande importância no sentido de ajudar o programador a entender possíveis melhorias que podem ser feitas no código objetivando a melhor performance possível.

Utilizando o gprof

Para perfilar o código C++ utilizando o gprof devem ser realizados os seguintes passos:

- Compilar o código com a flag -pg

Ex: `> g++ laplace.cxx -o laplace`

- Executar o executável gerado, o que gerará um arquivo gmon.out

Ex: `> ./laplace`

- Utilizar o gprof passando como argumento o arquivo gmon.out gerado no passo anterior - também é interessante armazenar a saída num arquivo de log (opcional)

Ex: `> gprof gmon.out [> log.txt]`

Relatório gprof

Flat profile:

Each sample counts as 0.01 seconds.

% cumulative	self	self	total			
time	seconds	seconds	calls	ms/call	ms/call	name
90.89	0.82	0.82	100	8.18	8.83	LaplaceSolver::timeStep(double)
7.25	0.88	0.07	24800400	0.00	0.00	SQR(double const&)
1.12	0.89	0.01	2	5.02	5.02	seconds()
0.56	0.90	0.01	2000	0.00	0.00	BC(double, double)
0.00	0.90	0.00	1	0.00	0.00	_GLOBAL__sub_I_ZN4GridC2Eii
0.00	0.90	0.00	1	0.00	0.00	__static_initialization_and_destruction_0(int, int)
0.00	0.90	0.00	1	0.00	0.00	LaplaceSolver::initialize()

0.00	0.90	0.00	1	0.00	883.28	LaplaceSolver::solve(int, double)
0.00	0.90	0.00	1	0.00	0.00	LaplaceSolver::LaplaceSolver(Grid*)
0.00	0.90	0.00	1	0.00	0.00	LaplaceSolver::~~LaplaceSolver()
0.00	0.90	0.00	1	0.00	5.02	Grid::setBCFunc(double (*)(double, double))
0.00	0.90	0.00	1	0.00	0.00	Grid::Grid(int, int)

% the percentage of the total running time of the
time program used by this function.

cumulative a running sum of the number of seconds accounted
seconds for by this function and those listed above it.

self the number of seconds accounted for by this
seconds function alone. This is the major sort for this
 listing.

calls the number of times this function was invoked, if
 this function is profiled, else blank.

self the average number of milliseconds spent in this
ms/call function per call, if this function is profiled,
 else blank.

total the average number of milliseconds spent in this
ms/call function and its descendents per call, if this
 function is profiled, else blank.

name the name of the function. This is the minor sort
 for this listing. The index shows the location of
 the function in the gprof listing. If the index is
 in parenthesis it shows where it would appear in
 the gprof listing if it were to be printed.

Copyright (C) 2012-2021 Free Software Foundation, Inc.

Copying and distribution of this file, with or without modification,
are permitted in any medium without royalty provided the copyright
notice and this notice are preserved.

Call graph (explanation follows)

granularity: each sample hit covers 2 byte(s) for 1.11% of 0.90 seconds

index	% time	self	children	called	name
					<spontaneous>
[1]	100.0	0.00	0.90		main [1]
		0.00	0.88	1/1	LaplaceSolver::solve(int, double) [3]
		0.01	0.00	2/2	seconds() [5]
		0.00	0.01	1/1	Grid::setBCFunc(double (*)(double, double)) [7]
		0.00	0.00	1/1	Grid::Grid(int, int) [19]
		0.00	0.00	1/1	LaplaceSolver::LaplaceSolver(Grid*) [17]

	0.00	0.00	1/1	LaplaceSolver::~LaplaceSolver() [18]
<hr/>				
	0.82	0.07	100/100	LaplaceSolver::solve(int, double) [3]
[2]	98.3	0.82	0.07 100	LaplaceSolver::timeStep(double) [2]
	0.07	0.00	24800400/24800400	SQR(double const&) [4]
<hr/>				
	0.00	0.88	1/1	main [1]
[3]	98.3	0.00	0.88 1	LaplaceSolver::solve(int, double) [3]
	0.82	0.07	100/100	LaplaceSolver::timeStep(double) [2]
<hr/>				
	0.07	0.00	24800400/24800400	LaplaceSolver::timeStep(double) [2]
[4]	7.3	0.07	0.00 24800400	SQR(double const&) [4]
<hr/>				
	0.01	0.00	2/2	main [1]
[5]	1.1	0.01	0.00 2	seconds() [5]
<hr/>				
	0.01	0.00	2000/2000	Grid::setBCFunc(double (*)(double, double)) [7]
[6]	0.6	0.01	0.00 2000	BC(double, double) [6]
<hr/>				
	0.00	0.01	1/1	main [1]
[7]	0.6	0.00	0.01 1	Grid::setBCFunc(double (*)(double, double)) [7]
	0.01	0.00	2000/2000	BC(double, double) [6]
<hr/>				
	0.00	0.00	1/1	__libc_csu_init [24]
[14]	0.0	0.00	0.00 1	__GLOBAL__sub_I_ZN4GridC2Eii [14]
	0.00	0.00	1/1	__static_initialization_and_destruction_0(int, int) [15]
<hr/>				
	0.00	0.00	1/1	__GLOBAL__sub_I_ZN4GridC2Eii [14]
[15]	0.0	0.00	0.00 1	__static_initialization_and_destruction_0(int, int) [15]
<hr/>				
	0.00	0.00	1/1	LaplaceSolver::LaplaceSolver(Grid*) [17]
[16]	0.0	0.00	0.00 1	LaplaceSolver::initialize() [16]
<hr/>				
	0.00	0.00	1/1	main [1]
[17]	0.0	0.00	0.00 1	LaplaceSolver::LaplaceSolver(Grid*) [17]
	0.00	0.00	1/1	LaplaceSolver::initialize() [16]
<hr/>				
	0.00	0.00	1/1	main [1]
[18]	0.0	0.00	0.00 1	LaplaceSolver::~LaplaceSolver() [18]
<hr/>				
	0.00	0.00	1/1	main [1]
[19]	0.0	0.00	0.00 1	Grid::Grid(int, int) [19]
<hr/>				

This table describes the call tree of the program, and was sorted by the total amount of time spent in each function and its children.

Each entry in this table consists of several lines. The line with the index number at the left hand margin lists the current function.

The lines above it list the functions that called this function, and the lines below it list the functions this one called.

This line lists:

index	A unique number given to each element of the table. Index numbers are sorted numerically. The index number is printed next to every function name so it is easier to look up where the function is in the table.
% time	This is the percentage of the `total' time that was spent in this function and its children. Note that due to different viewpoints, functions excluded by options, etc, these numbers will NOT add up to 100%.
self	This is the total amount of time spent in this function.
children	This is the total amount of time propagated into this function by its children.
called	This is the number of times the function was called. If the function called itself recursively, the number only includes non-recursive calls, and is followed by a `+' and the number of recursive calls.
name	The name of the current function. The index number is printed after it. If the function is a member of a cycle, the cycle number is printed between the function's name and the index number.

For the function's parents, the fields have the following meanings:

self	This is the amount of time that was propagated directly from the function into this parent.
children	This is the amount of time that was propagated from the function's children into this parent.
called	This is the number of times this parent called the function `/' the total number of times the function was called. Recursive calls to the function are not included in the number after the `/'.
name	This is the name of the parent. The parent's index number is printed after it. If the parent is a member of a cycle, the cycle number is printed between the name and the index number.

If the parents of the function cannot be determined, the word `

For the function's children, the fields have the following meanings:

self	This is the amount of time that was propagated directly from the child into the function.
------	---

children	This is the amount of time that was propagated from the child's children to the function.
called	This is the number of times the function called this child '/' the total number of times the child was called. Recursive calls by the child are not listed in the number after the '/'.
name	This is the name of the child. The child's index number is printed after it. If the child is a member of a cycle, the cycle number is printed between the name and the index number.

If there are any cycles (circles) in the call graph, there is an entry for the cycle-as-a-whole. This entry shows who called the cycle (as parents) and the members of the cycle (as children.) The '+' recursive calls entry shows the number of function calls that were internal to the cycle, and the calls entry for each member shows, for that member, how many times it was called from other members of the cycle.

Copyright (C) 2012-2021 Free Software Foundation, Inc.

Copying and distribution of this file, with or without modification, are permitted in any medium without royalty provided the copyright notice and this notice are preserved.

Index by function name

```
[14] _GLOBAL__sub_I__ZN4GridC2Eii [5] seconds() [17] LaplaceSolver::LaplaceSolver(Grid*)
[6] BC(double, double) [16] LaplaceSolver::initialize() [18] LaplaceSolver::~~LaplaceSolver()
[4] SQR(double const&) [3] LaplaceSolver::solve(int, double) [7] Grid::setBCFunc(double
*)(double, double))
[15] __static_initialization_and_destruction_0(int, int) [2] LaplaceSolver::timeStep(double) [19]
Grid::Grid(int, int)
```

Com o log gerado, observamos que o tempo total de execução foi de 0.90 s.

Considerando isso, temos um grande hotspot na função

LaplaceSolver::timeStep(double) e hotspots menores nas funções

SQR(double const&) - com um número muito alto de chamadas - e ***seconds()***.

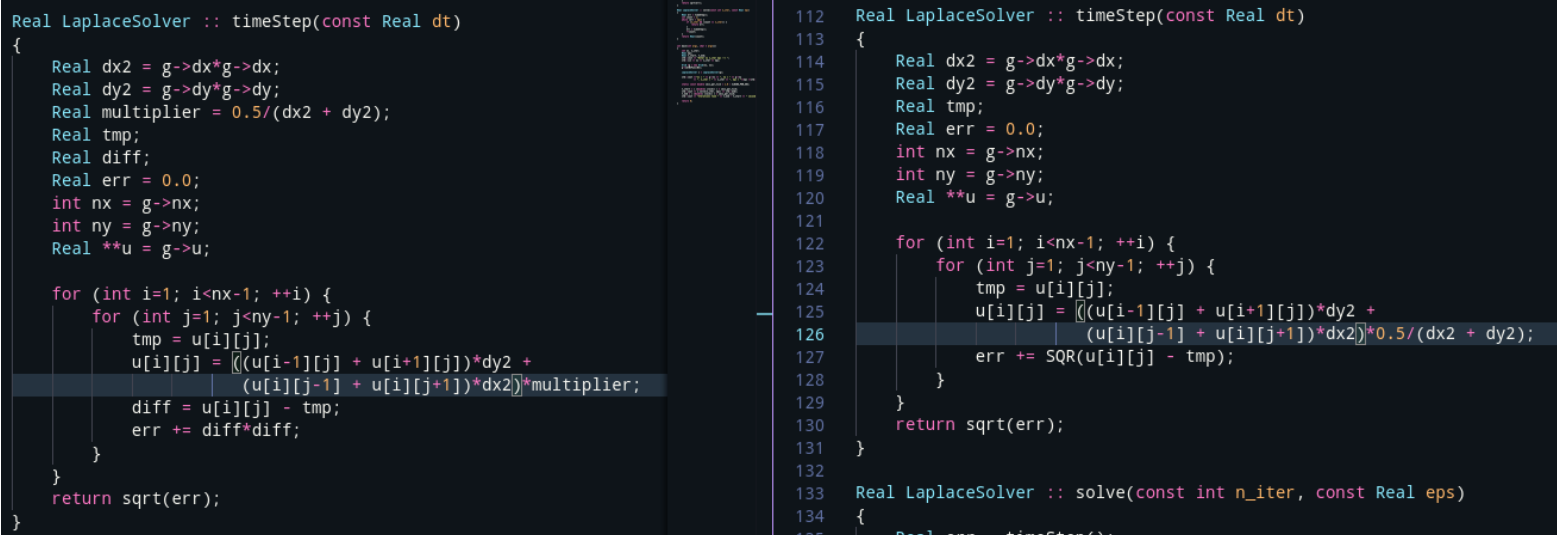
Observando o código, percebemos que este poderia sofrer algumas alterações seguindo as boas práticas de HPC. O código, no geral, segue estas diretrizes de forma parcial.

Alterações Realizadas

Visando a melhoria da performance do programa, foram feitas 3 alterações, cada uma delas atacando um dos hotspots mencionados anteriormente.

- **LaplaceSolver::timeStep(double):**

Nesta função, percebi que havia uma parte da equação executada dentro do loop que não precisava ser computada tantas vezes, dado que seu valor não mudava a cada iteração. Assim, foi feita a seguinte alteração:



The image shows a side-by-side comparison of two versions of the `LaplaceSolver::timeStep` function. The left version is the modified code, and the right version is the original code. In the modified version, the calculation of `0.5/(dx2 + dy2)` is moved out of the nested loops and stored in a variable named `multiplier`. This change is highlighted with a blue background in the original image. The original code calculates this value inside the loops, which is inefficient as it is recalculated for every iteration.

```
Real LaplaceSolver :: timeStep(const Real dt)
{
    Real dx2 = g->dx*g->dx;
    Real dy2 = g->dy*g->dy;
    Real multiplier = 0.5/(dx2 + dy2);
    Real tmp;
    Real diff;
    Real err = 0.0;
    int nx = g->nx;
    int ny = g->ny;
    Real **u = g->u;

    for (int i=1; i<nx-1; ++i) {
        for (int j=1; j<ny-1; ++j) {
            tmp = u[i][j];
            u[i][j] = ((u[i-1][j] + u[i+1][j])*dy2 +
                      (u[i][j-1] + u[i][j+1])*dx2)*multiplier;
            diff = u[i][j] - tmp;
            err += diff*diff;
        }
    }
    return sqrt(err);
}
```

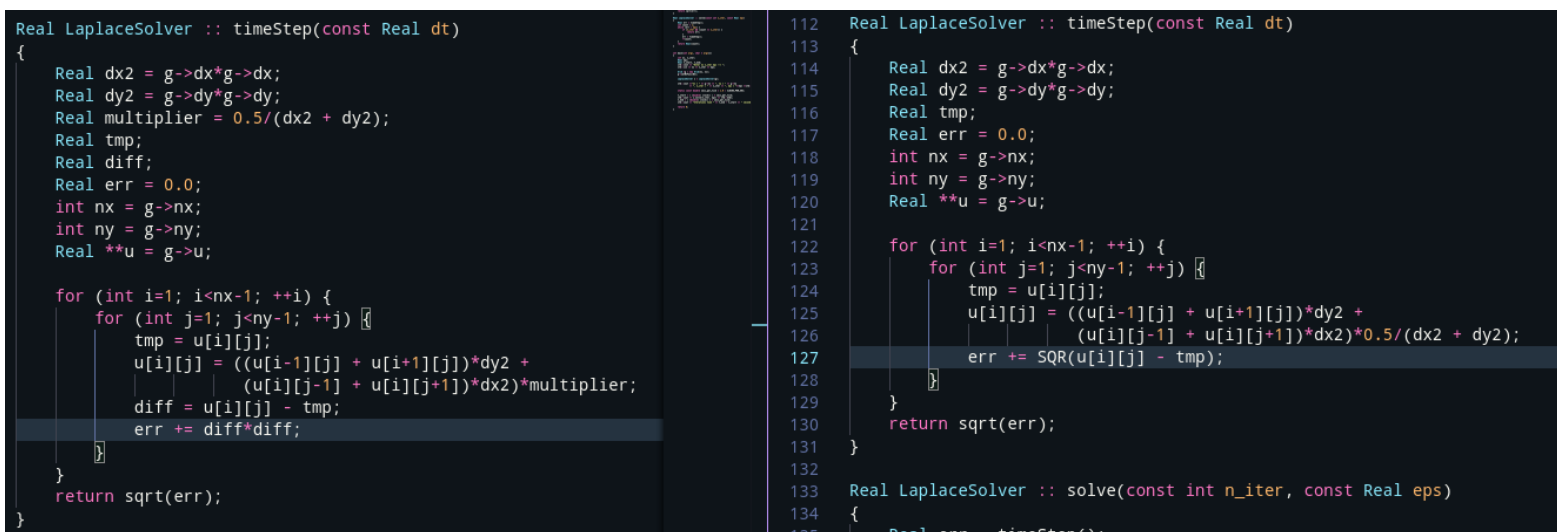
```
Real LaplaceSolver :: timeStep(const Real dt)
{
    Real dx2 = g->dx*g->dx;
    Real dy2 = g->dy*g->dy;
    Real tmp;
    Real err = 0.0;
    int nx = g->nx;
    int ny = g->ny;
    Real **u = g->u;

    for (int i=1; i<nx-1; ++i) {
        for (int j=1; j<ny-1; ++j) {
            tmp = u[i][j];
            u[i][j] = ((u[i-1][j] + u[i+1][j])*dy2 +
                      (u[i][j-1] + u[i][j+1])*dx2)*0.5/(dx2 + dy2);
            err += SQR(u[i][j] - tmp);
        }
    }
    return sqrt(err);
}
```

O arquivo à esquerda é a versão modificada, enquanto o da direita é a versão original. Aqui, simplesmente isolamos o $0.5(dx2 + dy2)$ como uma variável *multiplier* antes do loop, de forma a reaproveitar esta computação, melhorando a performance.

- **SQR(double const&)**

Esta função - como percebido no relatório do gprof - possui 24800400 chamadas, assim, a solução a ser utilizada foi retirar esta função do código e escrevê-la manualmente dentro do loop, único lugar onde ela é utilizada.



The image shows a side-by-side comparison of two versions of the `LaplaceSolver::timeStep` function. The left version is the modified code, and the right version is the original code. In the modified version, the `SQR` function is replaced by a direct calculation of `(u[i][j] - tmp) * (u[i][j] - tmp)` inside the loops. This change is highlighted with a blue background in the original image. The original code uses the `SQR` function, which is called many times, leading to performance issues.

```
Real LaplaceSolver :: timeStep(const Real dt)
{
    Real dx2 = g->dx*g->dx;
    Real dy2 = g->dy*g->dy;
    Real multiplier = 0.5/(dx2 + dy2);
    Real tmp;
    Real diff;
    Real err = 0.0;
    int nx = g->nx;
    int ny = g->ny;
    Real **u = g->u;

    for (int i=1; i<nx-1; ++i) {
        for (int j=1; j<ny-1; ++j) {
            tmp = u[i][j];
            u[i][j] = ((u[i-1][j] + u[i+1][j])*dy2 +
                      (u[i][j-1] + u[i][j+1])*dx2)*multiplier;
            diff = u[i][j] - tmp;
            err += diff*diff;
        }
    }
    return sqrt(err);
}
```

```
Real LaplaceSolver :: timeStep(const Real dt)
{
    Real dx2 = g->dx*g->dx;
    Real dy2 = g->dy*g->dy;
    Real tmp;
    Real err = 0.0;
    int nx = g->nx;
    int ny = g->ny;
    Real **u = g->u;

    for (int i=1; i<nx-1; ++i) {
        for (int j=1; j<ny-1; ++j) {
            tmp = u[i][j];
            u[i][j] = ((u[i-1][j] + u[i+1][j])*dy2 +
                      (u[i][j-1] + u[i][j+1])*dx2)*0.5/(dx2 + dy2);
            err += SQR(u[i][j] - tmp);
        }
    }
    return sqrt(err);
}
```

```
inline Real SQR(const Real &x)
{
    return (x*x);
}
```

- **seconds()**

Aqui, a alteração é análoga a feita com a função **SQR**, porém, como esta era uma função bem menos impactante para a performance - talvez inclusive nem possa ser considerada um hotspot - sua alteração não traz benefícios tão grandes assim como as outras, porém estes ainda sim são existentes.

Observando a função:

```
inline double seconds(void)
{
    static const double secs_per_tick = 1.0 / CLOCKS_PER_SEC;
    return ( (double) clock() ) * secs_per_tick;
}
```

Percebemos que a cada chamada desta, teremos a criação da const `secs_per_tick`, o que parece ser desnecessário. Assim, vamos retirar o comportamento desta função e escrevê-la manualmente onde ela é chamada:

```
int main(int argc, char * argv[])
{
    int nx, n_iter;
    Real eps;
    Real t_start, t_end;
    std::cout << "Enter nx n_iter eps --> ";
    std::cin >> nx >> n_iter >> eps;

    Grid *g = new Grid(nx, nx);
    g->setBCFunc(BC);

    LaplaceSolver s = LaplaceSolver(g);

    std::cout << "nx = " << g->nx << ", ny = " << g->ny
    << ", n_iter = " << n_iter << ", eps = " << eps << std::endl;

    static const double secs_per_tick = 1.0 / CLOCKS_PER_SEC;

    t_start = ( (double) clock() ) * secs_per_tick;
    std::cout << s.solve(n_iter, eps) << std::endl;
    t_end = ( (double) clock() ) * secs_per_tick;
    std::cout << "Iterations took " << t_end - t_start << " seconds.\n";

    return 0;
}

148 int main(int argc, char * argv[])
149 {
150     int nx, n_iter;
151     Real eps;
152     Real t_start, t_end;
153     std::cout << "Enter nx n_iter eps --> ";
154     std::cin >> nx >> n_iter >> eps;
155
156     Grid *g = new Grid(nx, nx);
157     g->setBCFunc(BC);
158
159     LaplaceSolver s = LaplaceSolver(g);
160
161     std::cout << "nx = " << g->nx << ", ny = " << g->ny
162     << ", n_iter = " << n_iter << ", eps = " << eps << std::endl;
163
164     t_start = seconds();
165     std::cout << s.solve(n_iter, eps) << std::endl;
166     t_end = seconds();
167     std::cout << "Iterations took " << t_end - t_start << " seconds.\n";
168
169     return 0;
170 }
171
```

Assim, passamos a criar a `secs_per_tick` apenas uma vez, otimizando ainda mais a performance do código.

Conclusão

Após realizar as alterações mencionadas anteriormente, rodamos o gprof novamente e obtivemos um ótimo resultado: o desempenho do programa melhorou consideravelmente.

A realização deste experimento deixou bem clara a importância de perfiladores - tais quais o gprof - como ferramenta de auxílio ao programador, principalmente aquele voltado à área de computação de alto desempenho. Este tipo de ferramenta se demonstrou muito útil, apontando quase que categoricamente pontos do código onde possivelmente caberia alguma melhoria

O novo tempo de execução foi de 52% do tempo de execução inicial, segue relatório abaixo.

Relatório:

Flat profile:

Each sample counts as 0.01 seconds.

	%	cumulative	self	self	total	
time	seconds	seconds	calls	ms/call	ms/call	name
100.38	0.47	0.47	100	4.72	4.72	LaplaceSolver::timeStep(double)
0.00	0.47	0.00	2000	0.00	0.00	BC(double, double)
0.00	0.47	0.00	1	0.00	0.00	_GLOBAL__sub_I_ZN4GridC2Eii
0.00	0.47	0.00	1	0.00	0.00	__static_initialization_and_destruction_0(int, int)
0.00	0.47	0.00	1	0.00	0.00	LaplaceSolver::initialize()
0.00	0.47	0.00	1	0.00	471.80	LaplaceSolver::solve(int, double)
0.00	0.47	0.00	1	0.00	0.00	LaplaceSolver::LaplaceSolver(Grid*)
0.00	0.47	0.00	1	0.00	0.00	LaplaceSolver::~~LaplaceSolver()
0.00	0.47	0.00	1	0.00	0.00	Grid::setBCFunc(double (*)(double, double))
0.00	0.47	0.00	1	0.00	0.00	Grid::Grid(int, int)

% the percentage of the total running time of the
time program used by this function.

cumulative a running sum of the number of seconds accounted
seconds for by this function and those listed above it.

self the number of seconds accounted for by this
seconds function alone. This is the major sort for this
 listing.

calls the number of times this function was invoked, if
 this function is profiled, else blank.

self the average number of milliseconds spent in this
ms/call function per call, if this function is profiled,
 else blank.

total the average number of milliseconds spent in this
ms/call function and its descendents per call, if this
 function is profiled, else blank.

name the name of the function. This is the minor sort
 for this listing. The index shows the location of
 the function in the gprof listing. If the index is
 in parenthesis it shows where it would appear in
 the gprof listing if it were to be printed.

Copyright (C) 2012-2021 Free Software Foundation, Inc.

Copying and distribution of this file, with or without modification,
are permitted in any medium without royalty provided the copyright
notice and this notice are preserved.

Call graph (explanation follows)

granularity: each sample hit covers 2 byte(s) for 2.12% of 0.47 seconds

index	% time	self	children	called	name
<spontaneous>					
[1]	100.0	0.00	0.47		main [1]
		0.00	0.47	1/1	LaplaceSolver::solve(int, double) [3]
		0.00	0.00	1/1	Grid::Grid(int, int) [17]
		0.00	0.00	1/1	Grid::setBCFunc(double (*)(double, double)) [16]
		0.00	0.00	1/1	LaplaceSolver::LaplaceSolver(Grid*) [14]
		0.00	0.00	1/1	LaplaceSolver::~~LaplaceSolver() [15]

		0.47	0.00	100/100	LaplaceSolver::solve(int, double) [3]
[2]	100.0	0.47	0.00	100	LaplaceSolver::timeStep(double) [2]

		0.00	0.47	1/1	main [1]
[3]	100.0	0.00	0.47	1	LaplaceSolver::solve(int, double) [3]
		0.47	0.00	100/100	LaplaceSolver::timeStep(double) [2]

		0.00	0.00	2000/2000	Grid::setBCFunc(double (*)(double, double)) [16]
[10]	0.0	0.00	0.00	2000	BC(double, double) [10]

		0.00	0.00	1/1	__libc_csu_init [22]
[11]	0.0	0.00	0.00	1	__GLOBAL__sub_I_ZN4GridC2Eii [11]
		0.00	0.00	1/1	__static_initialization_and_destruction_0(int, int) [12]

		0.00	0.00	1/1	__GLOBAL__sub_I_ZN4GridC2Eii [11]
[12]	0.0	0.00	0.00	1	__static_initialization_and_destruction_0(int, int) [12]

		0.00	0.00	1/1	LaplaceSolver::LaplaceSolver(Grid*) [14]
[13]	0.0	0.00	0.00	1	LaplaceSolver::initialize() [13]

		0.00	0.00	1/1	main [1]
[14]	0.0	0.00	0.00	1	LaplaceSolver::LaplaceSolver(Grid*) [14]
		0.00	0.00	1/1	LaplaceSolver::initialize() [13]

		0.00	0.00	1/1	main [1]
[15]	0.0	0.00	0.00	1	LaplaceSolver::~~LaplaceSolver() [15]

		0.00	0.00	1/1	main [1]
[16]	0.0	0.00	0.00	1	Grid::setBCFunc(double (*)(double, double)) [16]
		0.00	0.00	2000/2000	BC(double, double) [10]

		0.00	0.00	1/1	main [1]
[17]	0.0	0.00	0.00	1	Grid::Grid(int, int) [17]

This table describes the call tree of the program, and was sorted by the total amount of time spent in each function and its children.

Each entry in this table consists of several lines. The line with the

index number at the left hand margin lists the current function.

The lines above it list the functions that called this function,

and the lines below it list the functions this one called.

This line lists:

index A unique number given to each element of the table.
Index numbers are sorted numerically.
The index number is printed next to every function name so
it is easier to look up where the function is in the table.

% time This is the percentage of the 'total' time that was spent
in this function and its children. Note that due to
different viewpoints, functions excluded by options, etc,
these numbers will NOT add up to 100%.

self This is the total amount of time spent in this function.

children This is the total amount of time propagated into this
function by its children.

called This is the number of times the function was called.
If the function called itself recursively, the number
only includes non-recursive calls, and is followed by
a '+' and the number of recursive calls.

name The name of the current function. The index number is
printed after it. If the function is a member of a
cycle, the cycle number is printed between the
function's name and the index number.

For the function's parents, the fields have the following meanings:

self This is the amount of time that was propagated directly
from the function into this parent.

children This is the amount of time that was propagated from
the function's children into this parent.

called This is the number of times this parent called the
function '/' the total number of times the function
was called. Recursive calls to the function are not
included in the number after the '/'.
/

name This is the name of the parent. The parent's index
number is printed after it. If the parent is a
member of a cycle, the cycle number is printed between
the name and the index number.

If the parents of the function cannot be determined, the word
'<spontaneous>' is printed in the 'name' field, and all the other
fields are blank.

For the function's children, the fields have the following meanings:

self	This is the amount of time that was propagated directly from the child into the function.
children	This is the amount of time that was propagated from the child's children to the function.
called	This is the number of times the function called this child '/' the total number of times the child was called. Recursive calls by the child are not listed in the number after the '/'.
name	This is the name of the child. The child's index number is printed after it. If the child is a member of a cycle, the cycle number is printed between the name and the index number.

If there are any cycles (circles) in the call graph, there is an entry for the cycle-as-a-whole. This entry shows who called the cycle (as parents) and the members of the cycle (as children.) The '+' recursive calls entry shows the number of function calls that were internal to the cycle, and the calls entry for each member shows, for that member, how many times it was called from other members of the cycle.

Copyright (C) 2012-2021 Free Software Foundation, Inc.

Copying and distribution of this file, with or without modification, are permitted in any medium without royalty provided the copyright notice and this notice are preserved.

Index by function name

```
[11] _GLOBAL__sub_I__ZN4GridC2Eii [3] LaplaceSolver::solve(int, double) [16]
Grid::setBCFunc(double (*)(double, double))
[10] BC(double, double) [2] LaplaceSolver::timeStep(double) [17] Grid::Grid(int, int)
[12] __static_initialization_and_destruction_0(int, int) [14] LaplaceSolver::LaplaceSolver(Grid*)
[13] LaplaceSolver::initialize() [15] LaplaceSolver::~LaplaceSolver()
```

Todo o código utilizado pode ser acessado aqui: [Github](#)